

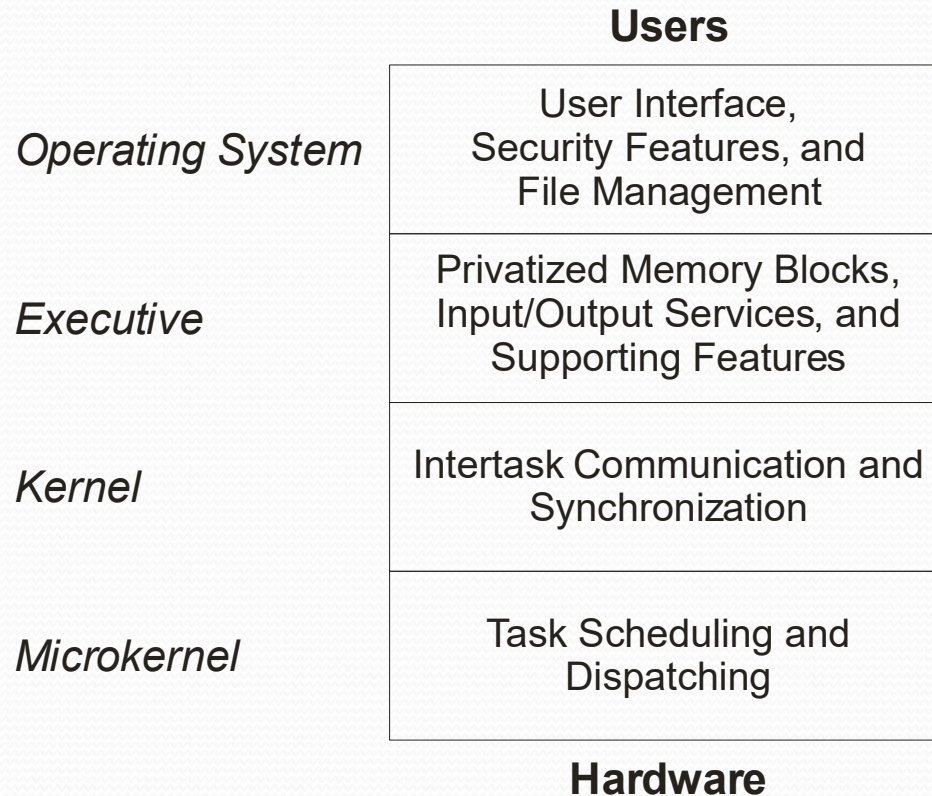
Real-Time & Embedded Systems

Agenda

- uC questions / wrapup
- RTOS (Chapter 3)



Operating systems taxonomy



The role of the kernel in operating systems. Moving up the taxonomy from the low-level microkernel to the full-featured operating system shows the additional functionality provided and also indicates the relative closeness to hardware versus human users.

Processes, and multiple threads

System

Process 1

Thread 1.1

Thread 1.2

Thread 1.3

Process 2

Thread 2.1

Thread 2.2

Thread 2.3

Thread 2.4

Process 3

Thread 3.1

Thread 3.2

Polled loop systems

- Simplest form of Real-time “kernel”.
- Used for fast reaction to single events
- Do not require interrupts.
- Simple while loop used to wait for an event signaled via DMA from a hardware device.
- Handler should clear event before servicing to catch bursts.
- These systems are easy to construct and analyze and thus to guarantee response times.

Polled loop systems


Hardware causes event via memory mapped I/O, DMA or a pin input

```
loop    {                               /* do forever    */
  if (packet_here)                       /* check flag    */
  {
    packet_here=0;                      /* reset flag    */
    process_data();                     /* process data  */
  }
}
```

Synchronized polled loops

- Polled loop response may be so fast that false events are induced by premature resetting of the flag.
- Switch bounce phenomenon in physical systems is a cause of such problems.
- Delay can be a no-op, loop, or interrupt (clock) driven delay.
- Increases response times-- but it's worth it.

Synchronized polled loops



Switch bounce phenomenon. The switch is closed at time t_0 , signaling the event, however, due to the ringing of the signal and the edge triggered logic several false events could be indicated at times t_1 , and t_2 .

Synchronized polled loops

```
loop    {                               /* do forever    */
        if(flag)                       /* check flag    */
        {
            pause(20);                 /* wait 20 ms    */
            flag=0;                    /* reset flag    */
            process_event();           /* process event */
        }
    }
```

wait for switch to
settle



Cyclic executives

- Simplest form of “multitasking.”
- Processes are run in round-robin fashion.
- Processes can be subdivided into “minor cycles” using FSM driven code.
- A small number of short processes will appear to be running concurrently.
- Easy to analyze for response times.
- Some people call periodic interrupt driven systems cyclic executives.

Cyclic executives

```
loop    {      /* do forever */
        check_for_keypressed();
        move.aliens();
        check_for_keypressed();
        check_for_collison()
        check_for_keypressed();
        update_screen();
    }
}
```

Cyclic executive for "Space Invaders."

Cyclic executives

- Not very flexible.
- Response times are very slow.
- All tasks run at the same “rate.”
- Tasks can be broken into states or phases to reduce response times and to create “minor” cycle.
- Can use FSM driver presented previously.
- Still does not require interrupts.

Coroutines

- Simplest form of “fairness scheduling.”
- Processes voluntarily exit at strategic points of execution
- Used with FSM.
- Inter-process communication through global variables.
- Sometimes called cooperative multitasking.
- Example: early versions of CICS.

Coroutines

```
void process_a(void)
{
loop {
    switch(state_a)
    {
        case 1:    phase_a1();
                   break;
        case 2:    phase_a2();
                   break;
        case 3:    phase_a3();
                   break;
        case 4:    phase_a4();
                   break;
        case 5:    phase_a5();
                   break;
    }
}
}
```

```
void process_b(void)
{
loop {
    switch(state_b)
    {
        case 1:    phase_b1();
                   break;
        case 2:    phase_b2();
                   break;
        case 3:    phase_b3();
                   break;
        case 4:    phase_b4();
                   break;
        case 5:    phase_b5();
                   break;
    }
}
}
```



Interrupt service routines

- Hardware interrupt: a signal generated by a peripheral device and sent to the CPU. The trigger is an electrical signal from an external device.
- Software interrupt: similar to the hardware interrupt, and it causes one code module to pass control to another. Trigger of is the execution of a machine language instruction.
- An exception is a software interrupt that is internal to the CPU and triggered by a software program's attempt to perform an unexpected or illegal operation.
- All three interrupts cause the CPU to transfer execution to a known location and then execute an interrupt service routine (ISR).



Interrupt service routines

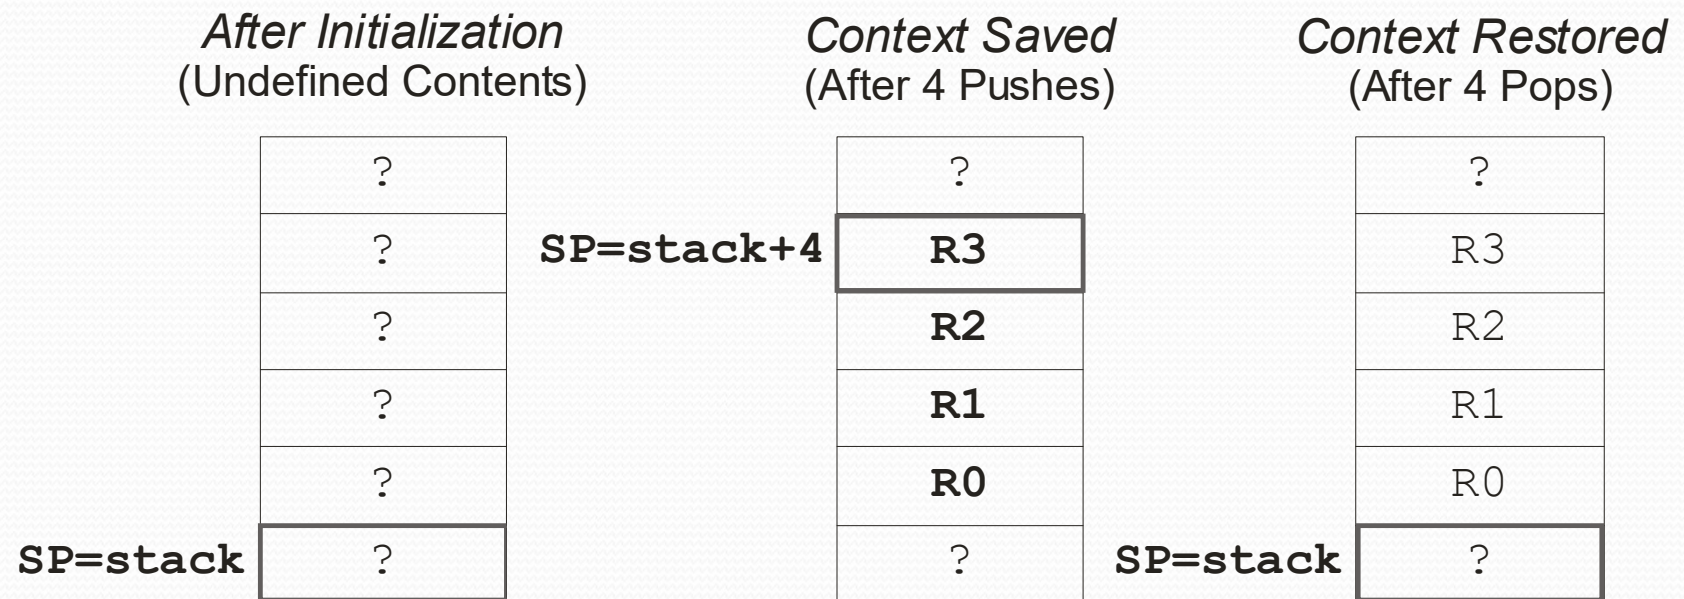
- Hardware interrupts are asynchronous in nature.
- Access to resources shared with an ISR is usually controlled by disabling interrupts in the application around any code that reads or writes to the resource.
- Synchronization mechanisms cannot be used in an ISR because the ISR cannot wait indefinitely for a resource to be available.
- If the ISR takes too long to process an interrupt, the external device may be kept waiting too long before its next interrupt is serviced.
- In all ISRs a snapshot of the machine—the context—must be preserved upon switching tasks so that it can be restored upon resuming the interrupted process.



Context switching

- Context switching is the process of saving and restoring sufficient information for a real-time task so that it can be resumed after being interrupted.
- The context is ordinarily saved to a stack data structure.
- Context switching time is a major contributor to response time and therefore must be minimized.
- The rule for saving context is simple: save the minimum amount of information necessary to safely restore any process after it has been interrupted.

Context switching



Context switching

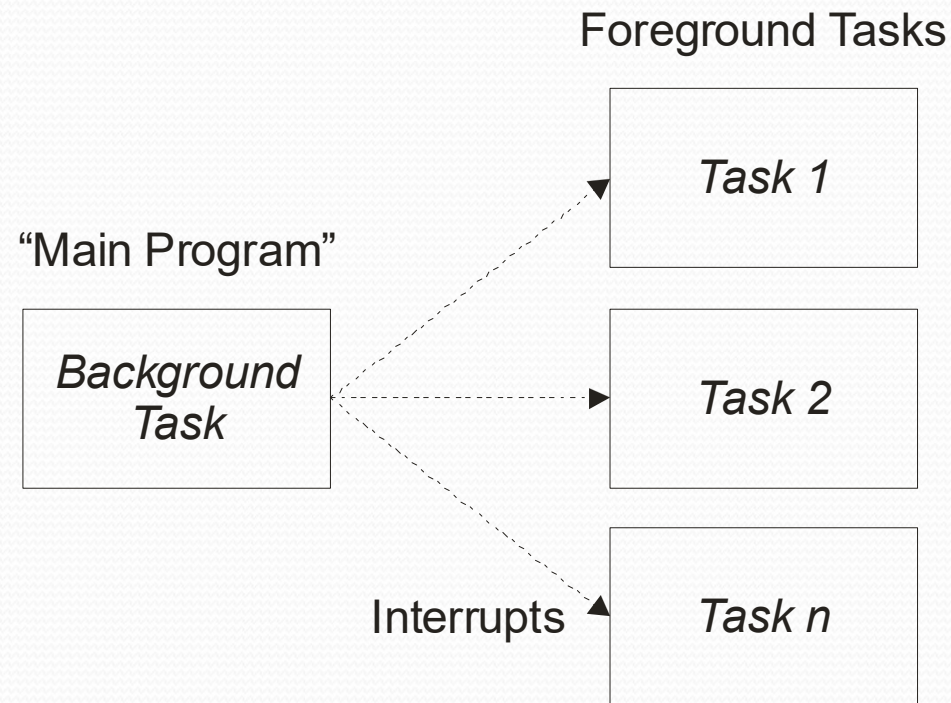
- Context usually includes:
 - contents of general registers
 - contents of the program counter
 - contents of coprocessor registers (if present)
 - memory page register
 - images of memory-mapped I/O locations (mirror images)
- Interrupts are disabled during context-switching. Sometimes a partial context switch is used to handle a burst of interrupts, to detect spurious interrupts, or to handle a time-overloaded condition.
- The stack model for context switching is used mostly in embedded systems where the number of real-time or interrupt-driven tasks is fixed.
- In the stack model, each interrupt handler is associated with a hardware interrupt and is invoked by the CPU, which vectors to the instruction stored at the appropriate interrupt-handler location.
- The context is then saved to a specially designated memory area that can be static, in the case of a single-interrupt system, or a stack, in the case of a multiple-interrupt system.



Preemptive priority systems

- Preemptive priority systems use preemption (prioritized interrupts). The priorities assigned to each interrupt are based on the urgency of the task associated with the interrupt.
- Prioritized interrupts can be either fixed priority or dynamic priority.
 - Fixed-priority systems are less flexible since the task priorities cannot be changed.
 - Dynamic-priority systems can allow the priority of tasks to be adjusted at run-time to meet changing process demands.
- Preemptive priority schemes can suffer from resource hogging by higher-priority tasks leading to a lack of available resources for lower-priority tasks. This is called starvation.
- Rate-monotonic systems are those fixed priority periodic real-time systems where the priorities are assigned so that the higher the execution frequency, the higher the priority.

Foreground/background systems



Foreground/background systems are the most common architecture for embedded applications. They involve a set of interrupt-driven or real-time processes called the foreground and a collection of non-interrupt driven processes called the background.



Background processing

- The background consist of all non-interrupt driven tasks. Includes anything that is not time critical.
- Foreground tasks are interrupt driven tasks and are time critical.
- Typical background tasks include:
 - low-priority self-testing
 - display updates
 - logging to printers, non-volatile memory
 - interfaces to slow devices



Initialization

- Initialization of the foreground/background system:
 - disable interrupts
 - set up interrupt vectors and stacks
 - perform self-test
 - perform system initialization
 - enable interrupts
- Initialization is actually the first part of the background process.
- It is important to disable interrupts because many systems come up with interrupts enabled while time is still needed to set things up.
- Initialization includes initializing the appropriate interrupt vector addresses, setting up stacks, and initializing any data, counters, arrays.
- Perform self-diagnostic tests before enabling any interrupts.
- Only then can real-time processing can begin.



Full-featured real-time operating systems

- The foreground/background solution can be extended into an operating system by adding additional functions such as network interfaces, complicated device drivers, and complex debugging tools.
- These types of systems are readily available as commercial products.
- Such systems rely on a complex operating system using round robin, preemptive priority, or a combination of both schemes to provide scheduling; the operating system represents the highest priority task, kernel, or supervisor.
- Commercial real-time operating systems are most often of this type. The task-control block model is most often used in these types of systems because the number of real-time tasks is indeterminate and dynamic.



The task control block model

- Each task is associated with a data structure – a task control block (TCB).
- TCB contains context. The system stores TCBs in one or more data structures, such as a linked list.
- TCB model can be used in round-robin, preemptive priority or combination systems, although it is generally associated with round-robin systems with a single clock.
- Is the most popular method for implementing commercial, full-featured, real-time operating systems.
- Also used in interactive on-line systems where tasks (associated with users) come and go.



The task control block model

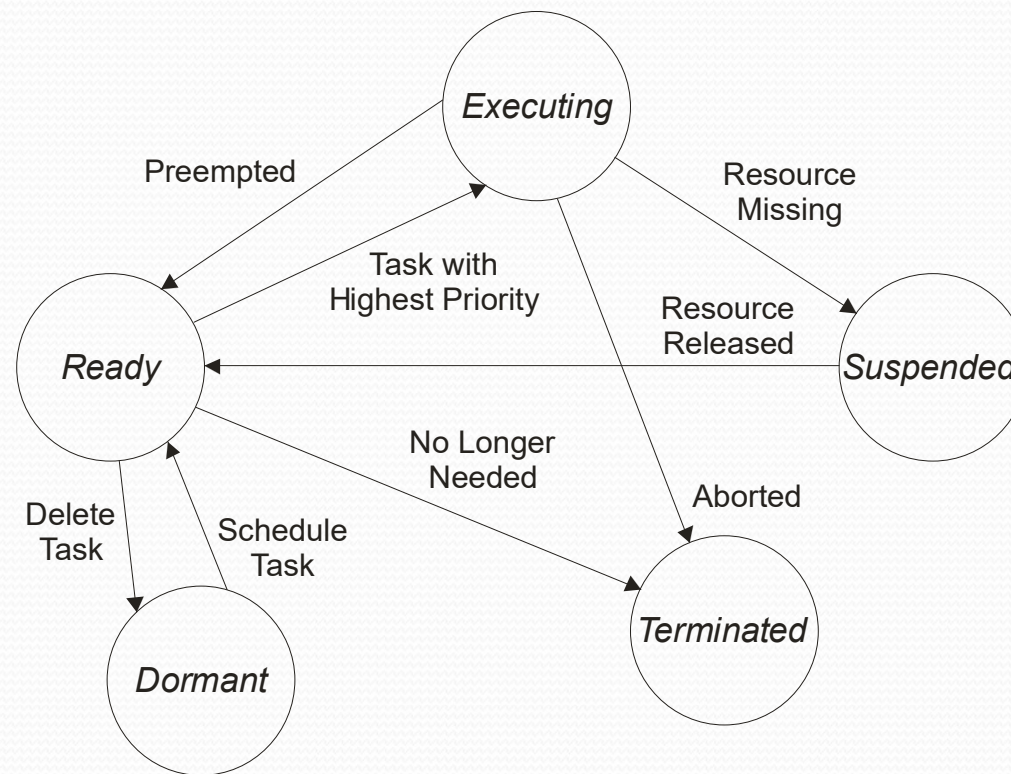
- The operating system manages TCBs by keeping track of the status of each task.
- A task can typically be in any one of the following states:
 - executing
 - ready
 - suspended (or blocked)
 - dormant (or sleeping)

Typical task control block

TCB

Task Identifier
Priority
Status
Work Registers
Program Counter
Status Register(s)
Stack Pointer
Pointer to Next TCB

The task control block model



A process state diagram as a partially defined finite state machine.



The task control block model

- Every hardware interrupt and every system level call (such as a request on a resource) invokes the real-time operating system.
- The operating system is responsible for maintaining a linked list containing the TCBs of all the ready tasks, and a second linked list of those in the suspended state.
- It also keeps a table of resources and a table of resource requests.



The task control block model

- In the TCB model tasks track their own resources.
- The TCB model is very flexible.
- The main drawback of the task-control block model is that when a large number of tasks are created, the overhead of the scheduler can become significant.

Buffering data

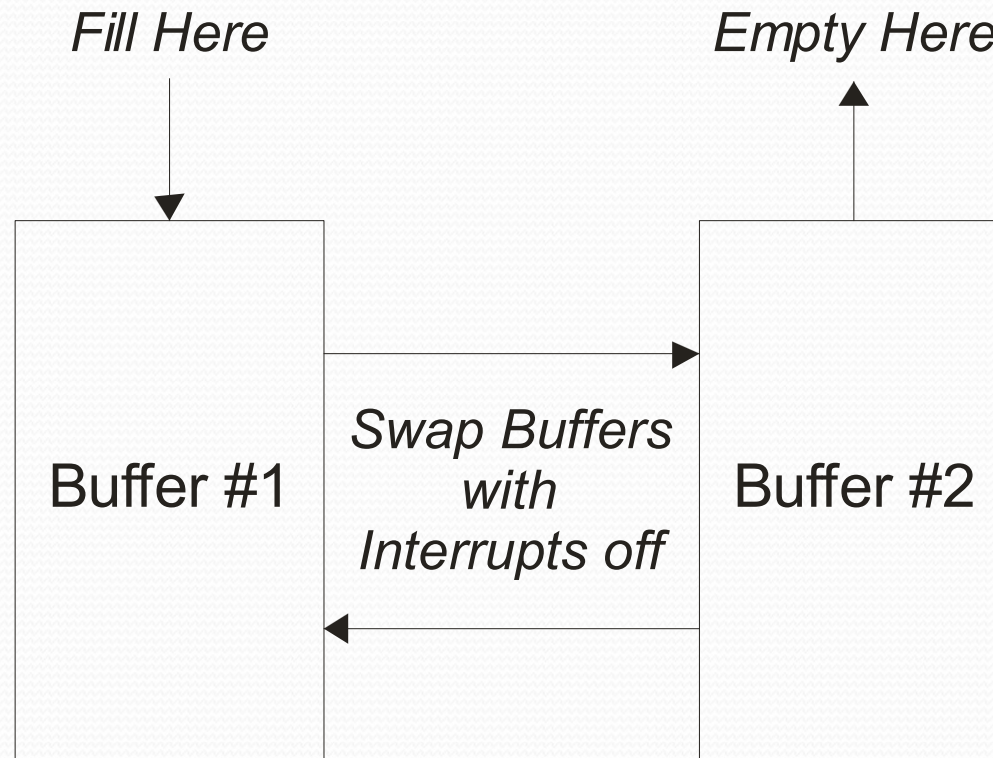
- Mechanisms are needed to pass data between tasks in a multitasking system when production and consumption rates are unequal.
- Global variables are simple and fast, but have collision potential.
- Example, one task may produce data at a constant 100 units per second, whereas another may consume these data at a rate less than 100 units per second.
 - Assuming that the production interval is finite (and relatively short), the slower consumption rate can be accommodated if the producer fills a storage buffer with the data.
 - The buffer holds the excess data until the consumer task can catch up.
 - The buffer can be a queue or other data structure, including an unorganized mass of variables.
 - If consumer task consumes this information faster than it can be produced, or if the consumer cannot keep up with the producer, problems occur.
- Selection of the appropriate size buffer and synchronization mechanisms is critical in reducing or eliminating these problems.



Time relative buffering

- Can use global variables for double buffering or ping-pong buffering.
- Used when time-relative (correlated) data need to be transferred between cycles of different rates, or when a full set of data is needed by one process but can only be supplied slowly by another process.
- Variant of the classic bounded buffer problem in which a block of memory is used as a repository for data produced by “writers” and consumed by “readers.”
- Further generalization is the readers and writers problem in which there are multiple readers and multiple writers of a shared resource

Time relative buffering




Double buffering configuration. Two identical buffers are filled and emptied by alternating tasks. Switching is accomplished either by a software pointer or hardware discrete.

Ring buffers

- A circular queue or ring buffer can be used to solve the problem of synchronizing multiple reader and writer tasks.
- Ring buffers are easier to manage than double buffers or queues when there are more than two readers or writers.
- Simultaneous input and output to the list are achieved by keeping head and tail indices.
- Data are loaded at the tail and read from the head.
- Can be used in conjunction with a counting or binary semaphore to control multiple requests for a single resource such as memory blocks, modems, and printers.

Ring buffers



A ring buffer. Processes write to the buffer at the head index and read data from the tail index. Data access is synchronized with a counting semaphore set to size of ring buffer.



Mailboxes

- A mutually agreed upon memory location that one or more tasks can use to pass data, or more generally for synchronization.
- Tasks rely on the kernel to allow them to write to the location via a `post` operation or to read from it via a `pend` operation.
- The difference between the `pend` operation and polling is that the pending task is suspended while waiting for data to appear -- eliminates the busy waiting condition.
- Mailboxes are available in most commercial RTOS.



Mailboxes

- The datum that is passed can be a flag used to protect a critical resource (called a key), a single piece of data, or a pointer to a data structure.
- In most implementations, when the key is taken from the mailbox, the mailbox is emptied.
- Since the key represents access to a critical resource, simultaneous access is precluded.
- Mailboxes are best implemented in systems based on the task control block model with a supervisor task.
- A table containing a list of tasks and needed resources is kept along with a second table containing a list of resources and their states.

Mailboxes

Task id #	Resource	Status
100	printer	has it
102	mailbox 1	has it
104	mailbox 1	pending

Task resource request table.

Resource	Status	Owner
printer 1	busy	100
mailbox 1	busy	102
mailbox 2	empty	none

Resource table used in conjunction with task resource request table.



Mailboxes

- When the supervisor is invoked by a system call or hardware interrupt, it checks the tables to see if some task is pending on a mailbox.
- If the key is available (key status is “full”), then than task must be restarted.
- If a task posts to a mailbox, then the operating system must ensure that the key is placed in the mailbox and its status updated to “full”.
- Other operations on the mailbox include the `accept` operation, which allows tasks to read the key if it is available, or immediately return an error code if the key is not available.
- In other implementations, the `pend` operation is equipped with a timeout, to prevent deadlocks.



Queues

- Some operating systems support `qpost`, `qpend`, and `qaccept` operations.
- In this case, the queue can be regarded as any array of mailboxes, and its implementation is facilitated through the same resource tables already discussed.
- Queues should not be used to pass arrays of data; pointers should be used instead.
- Queues are useful in implementing device servers where a pool of devices are involved:
 - The ring buffer holds requests for a device, and queues can be used at both the head and the tail to control access to the ring buffer.
 - This scheme is useful in the construction of device-controlling software.

Real-Time Systems Design and Analysis: Tools for the Practitioner

P. A. Laplante & S. J. Ovaska

Chapter 3
Part II

